

Exercise Sheet 12

Exercise 1: Building Neural Networks (20 + 20 P)

We consider the problem of learning decision functions in \mathbb{R}^2 where $x = (x_1, x_2)$ denotes the two-dimensional input vector. For this exercise, you only have access to neurons of the type

$$a_j = \sigma\left(b_j + \sum_i a_i w_{ij}\right)$$

where σ is the step function, i.e.

$$\sigma(t) = \begin{cases} 1 & \text{if } t > 0 \\ 0 & \text{if } t \leq 0 \end{cases}$$

(a) Construct a neural network that implements the decision boundary below

$$y = \begin{cases} 1 & \text{if } \max(x_1, x_2) > 2 \\ 0 & \text{if } \max(x_1, x_2) < 2 \end{cases}$$

Specifically, *draw* the neural network graph, and *indicate* for each neuron its weights and bias. (The exact behavior at the decision boundary does not need to be enforced.)

(b) Repeat the exercise for the decision function

$$y = \begin{cases} 1 & \text{if } \|x\|_1 > 2 \\ 0 & \text{if } \|x\|_1 < 2 \end{cases}$$

Exercise 2: Condition Number (20 + 10 P)

Consider a supervised dataset composed of inputs $\mathbf{x}_1, \dots, \mathbf{x}_N \in \mathbb{R}^d$ and respective targets $t_1, \dots, t_N \in \mathbb{R}$. Assume that $\frac{1}{N} \sum_{i=1}^N \mathbf{x}_i = \mathbf{0}$, i.e. the data is centered. Consider the homogeneous linear model $y = \mathbf{w}^\top \mathbf{x}$, with $\mathbf{w} \in \mathbb{R}^d$ a vector of parameters to be learned, and the regularized mean square objective:

$$\mathcal{E}(\mathbf{w}) = \alpha \|\mathbf{w}\|^2 + \frac{1}{N} \sum_{i=1}^N (\mathbf{w}^\top \mathbf{x}_i - t)^2$$

we would like to minimize.

(a) *Show* that the Hessian of the error function $\mathcal{E}(\mathbf{w})$ is given by the constant matrix:

$$H(\mathbf{w}) = 2(\Sigma + \alpha I)$$

where Σ is the covariance of the data.


(b) *Show* that the condition number associated to this Hessian matrix is given by

$$c = \frac{\lambda_1 + \alpha}{\lambda_d + \alpha}$$

where $\lambda_1, \dots, \lambda_d$ are the eigenvalues of the matrix Σ sorted in decreasing order.

Exercise 3: Backpropagation (30 P)

Consider some portion of a neural network given by:



(a) Assuming that we know the error gradient $(\partial E / \partial a_1, \partial E / \partial a_2)$, *compute* the error gradient $(\partial E / \partial z_1, \partial E / \partial z_2)$.