





WiSe 2024/25

Machine Learning 1/1-X



Lecture 14

# Product of Experts

# Outline

---

## Covered content:

- ▶ Products of Experts (PoE)
- ▶ Restricted Boltzmann Machine (RBM)
- ▶ Structure of an RBM
- ▶ RBM learning algorithms
- ▶ Application of RBMs

## Reference:

- ▶ Hinton, Geoffrey E. (2002). "Training Products of Experts by Minimizing Contrastive Divergence" (PDF). *Neural Computation*. 14 (8): 1771–1800

# Recap: Gaussian Mixture Models (1)

---

**Gaussian Mixture Model** represents a (multimodal) probability density function as a weighted sum of  $K$  Gaussian components

$$p(\mathbf{x}|\boldsymbol{\theta}) = \sum_{k=1}^K \alpha_k \cdot \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

where

- ▶  $K$  is the number of the individual components
- ▶  $\boldsymbol{\theta} = \{\alpha_1, \dots, \alpha_K, \boldsymbol{\mu}_1, \dots, \boldsymbol{\mu}_K, \boldsymbol{\Sigma}_1, \dots, \boldsymbol{\Sigma}_K\}$  are the model parameters.
- ▶  $\alpha_k$  with  $\alpha_k \geq 0$ ,  $\sum_k \alpha_k = 1$  are the mixing coefficients representing the weight of each Gaussian component. During sampling,  $\alpha_k$  gives the probability that  $\mathbf{x}$  has been generating by the  $k$ -th component.
- ▶  $\mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$  is the Gaussian distribution for component  $k$ , defined as

$$\mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) = \frac{1}{\sqrt{(2\pi)^d \det(\boldsymbol{\Sigma})}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu})\right)$$

# Recap: Gaussian Mixture Models (2)

---

- ▶ **Gaussian Mixture Model** are an example of latent variable models. For a latent variable  $y \in \{1, \dots, K\}$



$$p(x) = \sum_{k=1}^K p(x, y = k) = \sum_{k=1}^K \underbrace{p(y = k)}_{\alpha_k} \underbrace{p(x|y = k)}_{\mathcal{N}(x|\mu_k, \Sigma_k)}$$

where

- ▶ Given observed data points  $x_1, \dots, x_n$ , we can use the **Expectation Maximization** algorithm for **Maximum Likelihood Estimation** of the corresponding parameters  $\theta = \{\alpha_1, \dots, \alpha_K, \mu_1, \dots, \mu_K, \Sigma_1, \dots, \Sigma_K\}$ .
- ▶ Gaussian Mixture Models (GMMs) can approximate any smooth probability distribution to an arbitrary degree of accuracy.

# Recap: Gaussian Mixture Models (3)

---



# Product of Experts

---

Given  $K$  experts, each represented by an (unnormalized) probability density function  $g_k(x|\theta_k)$ , the combined model

$$p(x|\theta) = \frac{1}{\mathcal{Z}} \prod_{k=1}^K g_k(x|\theta_k)$$

is called **Product of Experts**, where

- ▶  $\theta = \{\theta_1, \dots, \theta_K\}$  are the model parameters
- ▶  $K$  is the number of the individual experts
- ▶ normalizing constant  $\mathcal{Z}$  is the partition function

$$\mathcal{Z} = \int \prod_{k=1}^K g_k(x|\theta_k) dx$$

# Example 1: Product of Gaussians (1)

For  $\mathbf{x} \in \mathbb{R}^d$ , consider the experts  $p_i(x_i) \sim \mathcal{N}(\mu_i, \sigma_i^2)$  to be univariate Gaussians specialized on a particular input dimension  $i \in \{1, \dots, d\}$  according to

$$p_i(\mathbf{x}|\mu_i, \sigma_i^2) = \frac{1}{\sqrt{2\pi\sigma_i^2}} \exp\left(-\frac{(x_i - \mu_i)^2}{2\sigma_i^2}\right)$$

The product of experts (see next slide for derivation) is given as:

$$\begin{aligned} p(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\sigma}^2) &= \prod_{i=1}^d p_i(\mathbf{x}|\mu_i, \sigma_i^2) \\ &= \frac{1}{\sqrt{(2\pi)^d \det(\boldsymbol{\Sigma})}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu})\right) \end{aligned}$$

with  $\boldsymbol{\Sigma} = \text{diag}(\sigma_1^2, \dots, \sigma_d^2)$ , i.e. a multivariate Gaussian distribution without correlation between features. Generally, any finite product of (multivariate) independent Gaussians is also a Gaussian with potentially correlated variables. Since the resulting distribution is a Gaussian, the product of experts (unlike mixture models) cannot approximate arbitrary smooth distribution. However, we gain efficiency due to the reduction in the number of model parameters.

## Example 1: Product of Gaussians (2)

---

For  $\boldsymbol{\mu} := (\mu_1, \dots, \mu_d)$ ,  $\Sigma := \text{diag}(\sigma_1^2, \dots, \sigma_d^2)$  and  $\Sigma^{-1} = \text{diag}(\frac{1}{\sigma_1^2}, \dots, \frac{1}{\sigma_d^2})$  we consider the following derivation of the exponential term:

$$\begin{aligned}\prod_{i=1}^d \exp\left(-\frac{(x_i - \mu_i)^2}{2\sigma_i^2}\right) &= \exp\left(-\frac{1}{2} \sum_{i=1}^d \frac{(x_i - \mu_i)^2}{\sigma_i^2}\right) \\ &= \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^\top \Sigma^{-1}(\mathbf{x} - \boldsymbol{\mu})\right),\end{aligned}$$

where  $\Sigma^{-1}(\mathbf{x} - \boldsymbol{\mu}) = (\frac{x_1 - \mu_1}{\sigma_1^2}, \dots, \frac{x_d - \mu_d}{\sigma_d^2})^\top$ .

Similarly, for the remaining terms we get:

$$\prod_{i=1}^d \frac{1}{\sqrt{2\pi\sigma_i^2}} = \frac{1}{\sqrt{(2\pi)^d \prod_{i=1}^d \sigma_i^2}} = \frac{1}{\sqrt{(2\pi)^d \det(\Sigma)}}$$

In particular,  $\mathcal{Z} = 1$ .



## Example 2: Product of Gaussian Mixtures

Let each expert be a mixture of univariate Gaussian distributions

$$p_i(\mathbf{x}|\theta_i) = \sum_{k=1}^K \alpha_{ik} \frac{1}{\sqrt{2\pi\sigma_{ik}^2}} \exp\left(-\frac{(x_i - \mu_{ik})^2}{2\sigma_{ik}^2}\right)$$

The product of experts (PoE) can be developed as:

$$\begin{aligned} p(\mathbf{x}|\theta) &= \frac{1}{\mathcal{Z}} \prod_{i=1}^d \sum_{k=1}^K \alpha_{ik} \frac{1}{\sqrt{2\pi\sigma_{ik}^2}} \exp\left(-\frac{(x_i - \mu_{ik})^2}{2\sigma_{ik}^2}\right) \\ &= \sum_{k_1=1}^K \cdots \sum_{k_d=1}^K \underbrace{\left( \frac{1}{\mathcal{Z}} \prod_{i=1}^d \alpha_{ik_i} \right)}_{\text{mix. coef.}} \underbrace{\left( \prod_{i=1}^d \frac{1}{\sqrt{2\pi\sigma_{ik_i}^2}} \exp\left(-\frac{(x_i - \mu_{ik_i})^2}{2\sigma_{ik_i}^2}\right) \right)}_{\text{multivar. Gaussian } \mathcal{N}(\boldsymbol{\mu}_{k_1, \dots, k_d}, \boldsymbol{\Sigma}_{k_1, \dots, k_d})} \end{aligned}$$

yielding a mixture of *exponentially* many ( $K^d$ ) multivariate Gaussians with uncorrelated variables, where  $\boldsymbol{\mu}_{k_1, \dots, k_d} = (\mu_{1k_1}, \dots, \mu_{dk_d})^\top$  and  $\boldsymbol{\Sigma}_{k_1, \dots, k_d} = \text{diag}(\sigma_{1k_1}^2, \dots, \sigma_{dk_d}^2)$ . Therefore, a PoE can encode (some!) mixture models by using fewer parameters. In general, PoE is less expressive.

## Example 3: Product of t-Student Distributions

---

Define  $H$  experts to be (a special case of) t-Student distributions in some projected space  $z = \mathbf{w}_i^\top \mathbf{x}$  with  $\|\mathbf{w}_i\| = 1$ :

$$p_i(\mathbf{x}|\theta_i) = \frac{1}{\alpha_i + (\mathbf{w}_i^\top \mathbf{x})^2}$$

The resulting product of experts

$$p(\mathbf{x}|\theta) = \frac{1}{Z} \prod_{j=1}^H \frac{1}{\alpha_j + (\mathbf{w}_j^\top \mathbf{x})^2}$$

produces a non-Gaussian multivariate distribution, which can be useful to model e.g. image or speech data. This PoE has connections to other analyses, e.g. independent component analysis (ICA).


# Mixture Models vs. Product of Experts

---

Aspect	Product of Experts (PoE)	Mixture Models
Combination Rule	Multiply distributions ( $\prod$ )	Add distributions ( $\sum$ )
Focus	Sharpens the distribution by focusing on intersections of experts.	Broadens the distribution by summing over components.
Interpretation	Experts act collaboratively to explain the data.	Components act competitively to explain the data.
Modeling	Captures complex interactions and dependencies.	Captures multi-modal distributions or clusters.
Normalization	Requires computing a potentially complex normalizer $\mathcal{Z}$ .	Weights $\pi_i$ handle normalization directly.

# The Restricted Boltzmann Machine

---



The restricted Boltzmann machine (RBM) is a joint probability model defined over input features  $\mathbf{x} \in \{0, 1\}^d$  **and** latent variables  $\mathbf{h} \in \{0, 1\}^H$ .

$$p(\mathbf{x}, \mathbf{h}|\boldsymbol{\theta}) = \frac{1}{Z} \exp \left( \sum_{i=1}^d \sum_{j=1}^H x_i w_{ij} h_j + \sum_{j=1}^H b_j h_j \right)$$

The parameter  $w_{ij}$  can be interpreted as the connection strength between input feature  $x_i$  and latent variable  $h_j$ . The larger  $w_{ij}$  the stronger  $x_i$  and  $h_j$  co-activate.

# The Restricted Boltzmann Machine (PoE View)

## Connection between RBM and PoE

The RBM, when marginalized over its hidden units, has the structure of a Product of Experts with  $g(\mathbf{x}|\theta_j) = (1 + \exp(\mathbf{w}_j^\top \mathbf{x} + b_j))$ .

*Proof:*

$$\begin{aligned} p(\mathbf{x}|\theta) &= \sum_{\mathbf{h} \in \{0,1\}^H} p(\mathbf{x}, \mathbf{h}|\theta) = \sum_{\mathbf{h} \in \{0,1\}^H} \frac{1}{\mathcal{Z}} \exp \left( \sum_{i=1}^d \sum_{j=1}^H x_i w_{ij} h_j + \sum_{j=1}^H b_j h_j \right) \\ &= \sum_{\mathbf{h} \in \{0,1\}^H} \frac{1}{\mathcal{Z}} \exp \left( \sum_{j=1}^H \left( \sum_{i=1}^d x_i w_{ij} \right) h_j + b_j h_j \right) \\ &= \sum_{\mathbf{h} \in \{0,1\}^H} \frac{1}{\mathcal{Z}} \exp \left( \sum_{j=1}^H (\mathbf{w}_j^\top \mathbf{x} + b_j) \cdot h_j \right) = \frac{1}{\mathcal{Z}} \sum_{\mathbf{h} \in \{0,1\}^H} \prod_{j=1}^H \exp((\mathbf{w}_j^\top \mathbf{x} + b_j) \cdot h_j) \\ &= \frac{1}{\mathcal{Z}} \prod_{j=1}^H \sum_{h_j \in \{0,1\}} \exp((\mathbf{w}_j^\top \mathbf{x} + b_j) \cdot h_j) \\ &= \frac{1}{\mathcal{Z}} \prod_{j=1}^H (1 + \exp(\mathbf{w}_j^\top \mathbf{x} + b_j)) \end{aligned}$$

# Interpreting the RBM Experts

---

The experts

$$g(\mathbf{x}|\boldsymbol{\theta}_j) = 1 + \exp(\mathbf{w}_j^\top \mathbf{x} + b_j)$$

forming the RBM implement two behaviors:

- ▶  $\mathbf{w}_j^\top \mathbf{x} + b_j > 0 : g(\mathbf{x}|\boldsymbol{\theta}_j) \gg 1$ : the example  $\mathbf{x}$  is in the area of competence of the expert and the latter speaks in favor of  $\mathbf{x}$ .
- ▶  $\mathbf{w}_j^\top \mathbf{x} + b_j < 0 : g(\mathbf{x}|\boldsymbol{\theta}_j) \approx 1$ : the example  $\mathbf{x}$  is outside the expert's area of competence, and the expert 'withdraws', i.e. multiplies by 1.

This double behavior is important to implement the nonlinearity.

# RBM as a Neural Network (1)


The product of expert (PoE) model

$$p(x|\theta) = \frac{1}{\mathcal{Z}} \prod_{j=1}^H (1 + \exp(\mathbf{w}_j^\top \mathbf{x} + b_j))$$

can be rewritten as:

$$\log p(x|\theta) = \underbrace{\sum_{j=1}^H \underbrace{\log(1 + \exp(\mathbf{w}_j^\top \mathbf{x} + b_j))}_{\text{softplus}}}_{f_\theta(x)} - \log \mathcal{Z}$$

where  $f_\theta(x)$  can be interpreted as a neural network. Here, the **softplus** operation implements the smooth version of the ReLU:  $x \mapsto \max(0, x)$ .




**Note:** The neural network can predict which examples  $x$  are more probable relative to other examples, but not the actual probability score, because  $\log \mathcal{Z}$  is difficult to compute.



# RBM as a Neural Network (2)

RBM trained on MNIST digits (only digits “3” with 100-125 foreground pixels).

2% least likely



2% most likely



## Observations:

- ▶ Digits with anomalies are predicted by the RBM to be less likely than clean digits.




# RBM as a Neural Network (3)

## Universal approximation

The RBM is an universal approximator of data distributions in the binary input space  $\{0, 1\}^d$ .


*"Proof" by construction:* Add a softplus neuron pointing to each corner of the hypercube. Scale the weight  $\mathbf{w}_j$  according to the probability of each corner, and choose the bias accordingly.




*Note:* This construction requires exponentially many neurons. In practice, each expert  $\theta_j$  captures more than a single corner of the hypercube (i.e. learn general features).

# RBM as a Neural Network (4)

RBM weights  $(\mathbf{w}_j)_{j=1}^H$



K-Means centroids



## Observation

- ▶ In the RBM, experts  $(\mathbf{w}_j)_j$  encode only part of the digit (e.g. a stroke). The expert therefore contributes to make all data containing that particular stroke more likely.
- ▶ The experts that the RBM extracts can be used for transfer learning, e.g. reused to predict different digits/characters.
- ▶ K-means centroids are much more localized, and consequently less transferable.

# Learning an RBM

## Recap: mixture model and EM bound (ELBO)

The mixture model can be optimized by finding the optimum of a lower-bound at each iteration

$$\log p(\mathcal{D}|\theta) \geq \sum_{n=1}^N \mathbb{E}_{q(z|x_n)} \left[ \log \frac{p(x_n, z|\theta)}{q(z|x_n)} \right] = \sum_{n=1}^N \sum_{k=1}^K \log \left( \frac{\alpha_k p(x_n|\theta_k)}{\beta_k} \right) \beta_k$$

where  $N$  is the number of data points, and  $(\alpha_k)_k$  and  $(\beta_k)_k$  are distributions over the  $K$  mixture elements.

**Question:** Can the same EM approach be used for Product of Experts?

$$\log p(\mathcal{D}|\theta) = \sum_{n=1}^N [f_{\theta}(x_n) - \log \mathcal{Z}]$$

**Answer:** No, the expression cannot be bounded in the same way as for the mixture model (it has a different structure).

# Gradient of the RBM

---

**Idea:** Although EM is not applicable, we can still compute the gradient of the log-likelihood and perform gradient descent.

$$\begin{aligned}\nabla_{\theta} \log p(\mathcal{D}|\theta) &= \nabla_{\theta} \sum_{n=1}^N [f_{\theta}(\mathbf{x}_n) - \log \mathcal{Z}] \\&= \nabla_{\theta} \sum_{n=1}^N \left[ f_{\theta}(\mathbf{x}_n) - \log \sum_{\mathbf{x} \in \{0,1\}^d} \exp(f_{\theta}(\mathbf{x})) \right] \\&= \sum_{n=1}^N \left[ \nabla_{\theta} f_{\theta}(\mathbf{x}_n) - \frac{\sum_{\mathbf{x} \in \{0,1\}^d} \exp(f_{\theta}(\mathbf{x})) \nabla_{\theta} f_{\theta}(\mathbf{x})}{\sum_{\mathbf{x} \in \{0,1\}^d} \exp(f_{\theta}(\mathbf{x}))} \right] \\&= \sum_{n=1}^N \nabla_{\theta} f_{\theta}(\mathbf{x}_n) - N \cdot \mathbb{E}_{\mathbf{x} \sim p(\mathbf{x}|\theta)} [\nabla_{\theta} f_{\theta}(\mathbf{x})]\end{aligned}$$

The gradient is a difference between a **data-dependent** and a **model-dependent** (i. e. data-independent) term.


# RBM Update Rule

Based on the gradient calculation above, we can build the update rule:

$$\theta \leftarrow \theta + \gamma \cdot \left( \frac{1}{N} \sum_{n=1}^N \nabla_{\theta} f_{\theta}(\mathbf{x}_n) - \mathbb{E}_{\mathbf{x} \sim p(\mathbf{x}|\theta)} \left[ \nabla_{\theta} f_{\theta}(\mathbf{x}) \right] \right)$$

**Interpretation:**

- ▶ The **first term** favours the model which makes the observed data more likely.
- ▶ The **second term** makes everything that the model considers likely, less likely.
- ▶ The training procedure terminates when we reach some equilibrium.



# Computation of the RBM update rule

---

$$\theta \leftarrow \theta + \gamma \cdot \left( \frac{1}{N} \sum_{n=1}^N \nabla_{\theta} f_{\theta}(\mathbf{x}_n) - \mathbb{E}_{\mathbf{x} \sim p(\mathbf{x}|\theta)} \left[ \nabla_{\theta} f_{\theta}(\mathbf{x}) \right] \right)$$


## Observations

- ▶ The **left term** is easy to compute (backprop in  $f_{\theta}(\mathbf{x})$ , averaged over all data points).
- ▶ The **right term** is more tricky, because  $p(\mathbf{x}|\theta)$  is defined over exponentially many states. An unbiased approximation of the expected gradient can be obtained by generating a sample  $\{\mathbf{x}_1, \dots, \mathbf{x}_m\}$  from the model distribution  $p(\mathbf{x}|\theta)$ .

**Question:** How do we sample from  $p(\mathbf{x}|\theta)$ ?

- ▶ *Idea:* Switch back to the 'classical' (i.e. non-PoE) view of the RBM, where we can use latent variables to ease sampling.

# Recap: 'Classical' View of the RBM



The RBM is a probability model defined over input features  $\mathbf{x} \in \{0, 1\}^d$  and **and** latent variables  $\mathbf{h} \in \{0, 1\}^H$ .

$$p(\mathbf{x}, \mathbf{h} | \boldsymbol{\theta}) = \frac{1}{Z} \exp \left( \sum_{i=1}^d \sum_{j=1}^H x_i w_{ij} h_j + \sum_{j=1}^H b_j h_j \right)$$

The parameter  $w_{ij}$  can be interpreted as the connection strength between input feature  $x_i$  and latent variable  $h_j$ . The larger  $w_{ij}$  the stronger  $x_i$  and  $h_j$  co-activate.

# Conditional Distributions in an RBM

---

Sampling  $p(\mathbf{x})$  and  $p(\mathbf{h})$  is hard, however, hidden variables  $\mathbf{h}$  can be sampled *easily* and *independently* when we condition on the state  $\mathbf{x}$ , and similarly for  $\mathbf{x}$  conditioned on  $\mathbf{h}$ . Let  $z_j = \sum_i x_i w_{ij} + b_j$ . We proceed as:

$$p(\mathbf{h}|\mathbf{x}, \theta) = \frac{\frac{1}{\mathcal{Z}} \exp\left(\sum_j z_j h_j\right)}{\sum_{\mathbf{h}} \frac{1}{\mathcal{Z}} \exp\left(\sum_j z_j h_j\right)} = \frac{\prod_j \exp(z_j h_j)}{\prod_j \sum_{h_j} \exp(z_j h_j)} = \prod_j \underbrace{\frac{\exp(z_j h_j)}{1 + \exp(z_j)}}_{p(h_j|\mathbf{x}, \theta)}$$

and we observe that, conditioned on  $\mathbf{x}$ , the latent variables  $(h_j)_j$  can be sampled *easily* (Bernoulli distributions) and *independently* (hidden variables in RBM are conditionally independent given  $\mathbf{x}$ ).

By symmetry of the RBM model, we get a similar result for the visible units conditioned on the latent variables, i.e.  $p(x_i|\mathbf{h}, \theta) = \exp(z_i h_i)/(1 + \exp(z_i))$  with  $z_i = \sum_j w_{ij} h_j$ .



# Block Gibbs Sampling

---

**Observation:** We know  $p(h_j | \mathbf{x}, \theta)$  and  $p(x_i | \mathbf{h}, \theta)$ . But how do we sample  $p(\mathbf{x} | \theta)$ , which we need to compute the RBM gradient?

**Block Gibbs Sampling:** Start with some random input  $\mathbf{x}^{(0)}$ , then sample alternately:

$$\begin{aligned}\mathbf{h}^{(0)} &\sim (p(h_j | \mathbf{x}^{(0)}, \theta))_{j=1}^H \\ \mathbf{x}^{(1)} &\sim (p(x_i | \mathbf{h}^{(0)}, \theta))_{i=1}^d \\ \mathbf{h}^{(1)} &\sim (p(h_j | \mathbf{x}^{(1)}, \theta))_{j=1}^H \\ \mathbf{x}^{(2)} &\sim (p(x_i | \mathbf{h}^{(1)}, \theta))_{i=1}^d \\ &\vdots\end{aligned}$$

The procedure guarantees that  $\mathbf{x}^{(\cdot)}$  converges in distribution to  $p(\mathbf{x} | \theta)$ , that is,

$$\lim_{n \rightarrow \infty} F_{X^{(n)}}(\mathbf{x}) = F_X(\mathbf{x})$$

where  $F_X(\mathbf{x}) = P(X_1 \leq x_1, \dots, X_d \leq x_d)$  is the cumulative distribution function.

# Fast Approximations

---

In practice, Gibbs sampling can take a long time to converge. Instead, we can use fast approximations of converged Gibbs sampling.

## Common approximation strategies:

- ▶ *Contrastive divergence (CD- $k$ )*: Start from an existing data point, and perform  $k$  steps of alternate Gibbs sampling.
- ▶ *Persistent contrastive divergence (PCD)*: Start from the Gibbs sample  $x$  in the previous iteration of gradient descent, and perform one step of Gibbs sampling.

# Application: RBM for Data Denoising

---

Suppose you receive an example  $\mathbf{x}_{\text{noisy}}$ , and would like to denoise it. A reconstruction of that digit can be obtained by mapping to the latent variables and projecting back:


1. Projection on latent variables

$$\mathbf{h} = \text{sigm}(W\mathbf{x}_{\text{noisy}} + \mathbf{b})$$

2. Backprojection on the input domain

$$\mathbf{x}_{\text{rec}} = \text{sigm}(W^T \mathbf{h})$$

Before denoising



After denoising



# Application: RBM for Representation Learning


The RBM model can be used as a neural network initialization to achieve lower generalization error on some classification task.

- **Step 1:** Create a layer of features based on the RBM learned parameters:

$$\phi(\mathbf{x}) = \begin{pmatrix} \text{sigm}(\mathbf{w}_1^\top \mathbf{x}) \\ \vdots \\ \text{sigm}(\mathbf{w}_H^\top \mathbf{x}) \end{pmatrix}$$

- **Step 2:** Add a top layer that maps  $\phi(\mathbf{x})$  to the class probabilities, and train the resulting neural network end-to-end using gradient descent.

## MNIST example:



Source: Erhan et al. (2010) Why Does Unsupervised Pre-training Help Deep Learning?

# Summary

---

- ▶ The Product of Experts is an unsupervised learning approach that is substantially different from Mixture Models.
- ▶ Product of experts such as the RBM are optimized by gradient descent (instead of EM).
- ▶ The RBM has several desirable properties: Simple gradient, block Gibbs sampler (quite fast).
- ▶ The RBM can be used for various tasks, e.g. probability modeling, data denoising, representation learning.